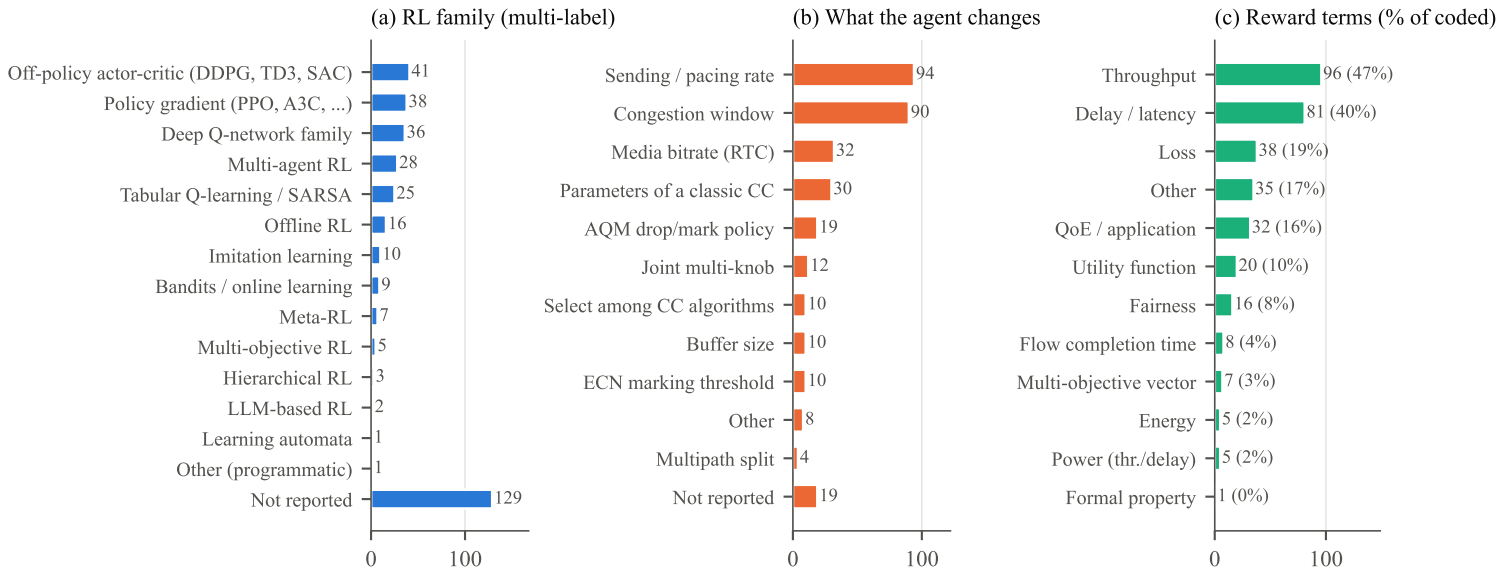
(a) RL family (multi-label)
Off-policy actor-critic (DDPG, TD3, SAC)
41
Policy gradient (PPO, A3C, ...)
38
Deep Q-network family
36
Multi-agent RL
28
Tabular Q-learning / SARSA
25
Offline RL
16
Imitation learning
10
Bandits / online learning
9
Meta-RL
7
Multi-objective RL
5
Hierarchical RL
3
LLM-based RL
2
Learning automata
1
Other (programmatic)
1
Not reported
129
0
100
(b) What the agent changes
Sending / pacing rate
94
Congestion window
90
Media bitrate (RTC)
32
Parameters of a classic CC
30
AQM drop/mark policy
19
Joint multi-knob
12
Select among CC algorithms
10
Buffer size
10
ECN marking threshold
10
Other
8
Multipath split
4
Not reported
19
0
100
(c) Reward terms (% of coded)
Throughput
96 (47%)
Delay / latency
81 (40%)
Loss
38 (19%)
Other
35 (17%)
QoE / application
32 (16%)
Utility function
20 (10%)
Fairness
16 (8%)
Flow completion time
8 (4%)
Multi-objective vector
7 (3%)
Energy
5 (2%)
Power (thr./delay)
5 (2%)
Formal property
1 (0%)
0
100